

MASTER RESEARCH INTERNSHIP IN COMPUTER SCIENCE

Machine learning, Information and Content

“Unsupervised Neural Word Alignment HMM”

Computer Science Laboratory for Mechanics and Engineering Sciences (LIMSI)

Author:
VU Trong-Bach

Supervisor:
Dr. François YVON

Contents



I. ISSUES INTRODUCTION

II. PROPOSED METHOD

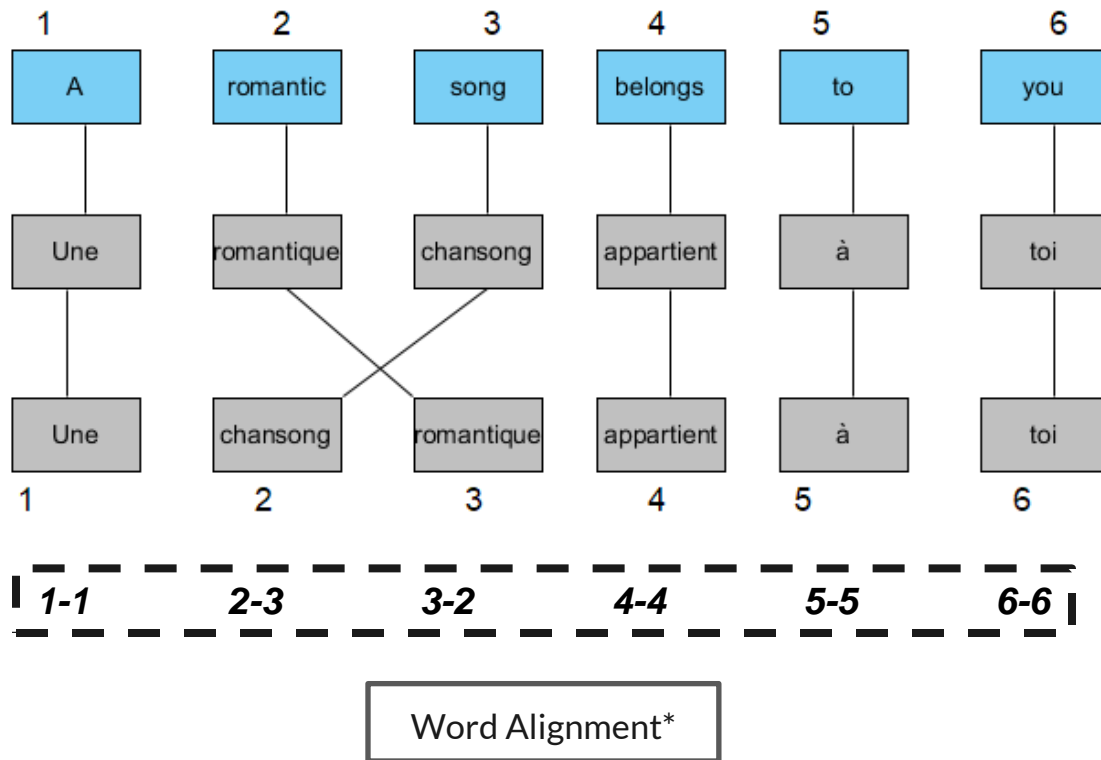

III. EXPERIMENTS & RESULTS

IV. CONCLUSION



I. ISSUES INTRODUCTION

I. ISSUES INTRODUCTION



I. ISSUES INTRODUCTION



$$f^J = \{f_1 \cdots f_j \cdots f_J\}$$

$$e^I = \{e_1 \cdots e_i \cdots e_I\}$$

$$e^I = \arg \max_{e^I} \{P(e^I | f^J)\}$$

$$= \arg \max_{e^I} \{P(e^I) \cdot P(f^J | e^I)\}$$




$$P(f^J | e^I) = \sum_{a^J} P(f^J, a^J | e^I)$$

$$= \sum_{a^J} \prod_{j=1}^J P(f_j, a_j | f_{j-1}, a_{j-1}, e^I)$$

$$= \sum_{a^J} \prod_{j=1}^J P(a_j | f_{j-1}, a_{j-1}, e^I) \cdot P(f_j | f_{j-1}, a_j, e^I)$$

First order dependence




$$P(f^J | e^J) = \sum_{a^J} \prod_{j=1}^J [p(a_j | a_{j-1}, I) \cdot p(f_j | e_{a_j})]$$

Hidden Markov Models


$$P(a_j | f_{j-1}, a_{j-1}, e^I) = p(a_j | a_{j-1}, I)$$

$$P(f_j | f_{j-1}, a_j, e^I) = p(f_j | e_{a_j})$$

I. ISSUES INTRODUCTION



Hidden Markov Models



Neural Network

$$P(f^J | e^J) = \sum_{a^J} \prod_{j=1}^J [\underbrace{p(a_j | a_{j-1}, I)}_{\text{Transition Model or Alignment Model}} \cdot \underbrace{p(f_j | e_{a_j})}_{\text{Emission Model or Translation Model}}]$$

**Transition Model
or Alignment Model**

**Emission Model
or Translation Model**

Neuralized



II. PROPOSED METHOD



**Transition Model
or Alignment Model**

**Emission Model
or Translation Model**

Neutralized

II. PROPOSED METHOD

Transition Model

$$p(a_j|a_{j-1}, I) \text{ or } p(i|i', I) = \frac{s(i - i')}{\sum_{i''=1}^I s(i'' - i')}$$

Using a
non-negative set

i'/i	0	1	2	...
0	s(0)	s(1)	s(2)	...
1	s(-1)	s(0)	s(1)	s(2)
2	s(-2)	s(-1)	s(0)	s(1)
...	...	s(-2)	s(-1)	s(0)

Empty word problem???

II. PROPOSED METHOD

Transition Model

$$p(i + I|i', I) = p_0 \cdot \delta(i, i')$$

$$p(i + I|i' + I, I) = p_0 \cdot \delta(i, i')$$

$$p(i|i' + I, I) = p(i|i', I)$$


- p_0 is the probability of a transition to the empty word
- $\delta(i, i') \begin{cases} 1 & \text{if } i = i' \\ 0 & \text{otherwise} \end{cases}$

by extending the HMM empty words e^{2I}

i'/i	0	1	2	...	0 (I)	1 (I+1)	2 (I+2)	...
0	$s(0)$	$s(1)$	$s(2)$...	$p_0 \cdot 1$	$p_0 \cdot 0$	$p_0 \cdot 0$...
1	$s(-1)$	$s(0)$	$s(1)$	$s(2)$	$p_0 \cdot 0$	$p_0 \cdot 1$	$p_0 \cdot 0$	$p_0 \cdot 0$
2	$s(-2)$	$s(-1)$	$s(0)$	$s(1)$	$p_0 \cdot 0$	$p_0 \cdot 0$	$p_0 \cdot 1$	$p_0 \cdot 0$
...	...	$s(-2)$	$s(-1)$	$s(0)$...	$p_0 \cdot 0$	$p_0 \cdot 0$	$p_0 \cdot 1$
0 (I)	$s(0)$	$s(1)$	$s(2)$...	$p_0 \cdot 1$	$p_0 \cdot 0$	$p_0 \cdot 0$...
1 (I+1)	$s(-1)$	$s(0)$	$s(1)$	$s(2)$	$p_0 \cdot 0$	$p_0 \cdot 1$	$p_0 \cdot 0$	$p_0 \cdot 0$
2 (I+2)	$s(-2)$	$s(-1)$	$s(0)$	$s(1)$	$p_0 \cdot 0$	$p_0 \cdot 0$	$p_0 \cdot 1$	$p_0 \cdot 0$
...	...	$s(-2)$	$s(-1)$	$s(0)$...	$p_0 \cdot 0$	$p_0 \cdot 0$	$p_0 \cdot 1$

II. PROPOSED METHOD

Emission Model



Unsupervised - How to update θ ???

II. PROPOSED METHOD

Update θ

Maximize the evidence


$$p(f|\theta) = \sum_e p(f, e|\theta)$$

➤ To estimate θ , we can use auxiliary function of EM algorithm

$$p(f|\theta) = E_{q(e)}[\ln p(f, e|\theta)] + H[q(e)] + KL(q(e)||p(f, e|\theta))$$

We choose:

- $q(e)$ to be posterior $p(e|f)$
- $H[q(e)]$ a constant \rightarrow dropped
- Setting KL divergence to zero



Only maximize $E_{p(e|f)}[\ln p(f, e|\theta)]$

Gradient:

$$J(\theta) = \sum_e p(e|f) \frac{\partial}{\partial \theta} \ln p(f, e|\theta)$$

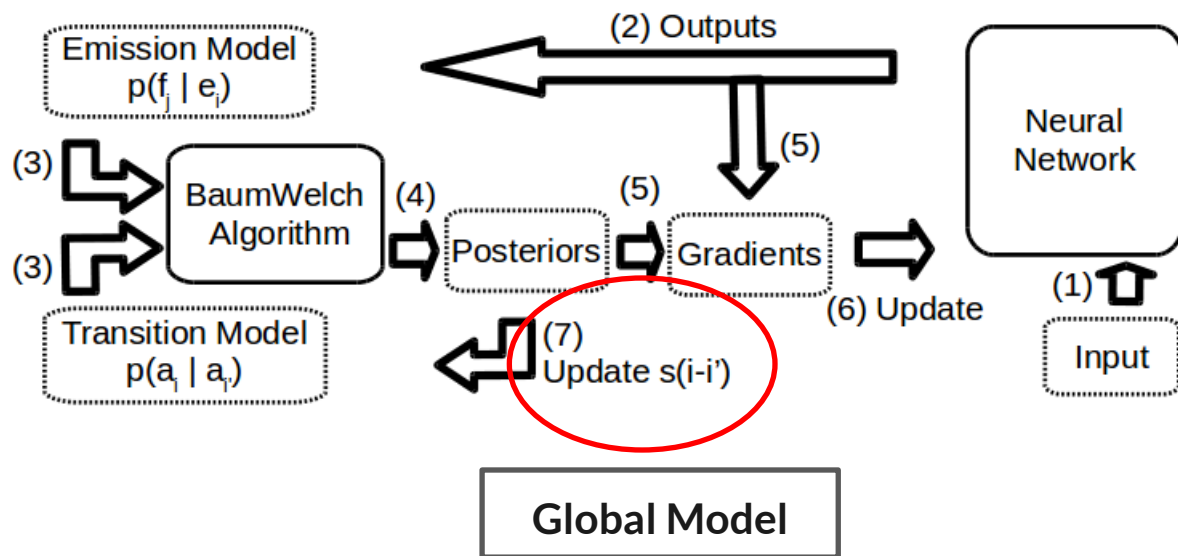
$$J(\theta) = \sum_j \sum_{a_j} p(a_j|f_j) \frac{\partial}{\partial \theta} \ln p(f_j|e_{a_j}, \theta)$$

How calculate posteriors ???

II. PROPOSED METHOD

$$p(a_j = i' | f_j, \theta) \propto \alpha_{i'}(j) \beta_{i'}(j)$$

$$p(a_j = i', a_{j+1} = i | f^J, \theta) \propto \alpha_{i'}(j) p(a_{j+1} = i | a_j = i') \times \beta_i(j+1) p(f_{j+1} | e_{a_j})$$



II. PROPOSED METHOD

Update non-negative set $s(i-i')$

$$p(a_j = i' | f_j, \theta) \propto \alpha_{i'}(j) \beta_{i'}(j)$$

$$p(a_j = i', a_{j+1} = i | f^J, \theta) \propto \alpha_{i'}(j) p(a_{j+1} = i | a_j = i') \times \beta_i(j+1) p(f_{j+1} | e_{a_j})$$

$$s(i, i') = \frac{\sum_n^N \sum_j^{J-1} p_n(a_j = i', a_{j+1} = i | f^J, \theta)}{\sum_n^N \sum_j^{J-1} p_n(a_j = i' | f_j, \theta)}$$

*N is the number of pair
sentences (f; e) in the corpus*



III. EXPERIMENTS

III. EXPERIMENTS


Evaluation Methodology

Viterbi Alignment

$$\begin{aligned}\hat{a}^J &= \arg \max_{a^J} p(f^J, a^J | e^{2I}) \\ &= \arg \max_{a^J} \left\{ \prod_{j=1}^J [p(a_j | a_{j-1}, 2I) \cdot p(f_j | e_{a_j}^{2I})] \right\} \\ &= \left[\arg \max_{a_j} \{p(a_j | a_{j-1}, 2I) \cdot p(f_j | e_{a_j}^{2I})\} \right]_{j=1}^J\end{aligned}$$

$$AER = 1 - \frac{(|A \cap S| + |A \cap P|)}{(|A| + |S|)}$$

The best AER = 0.0,
 where S (sure alignments), P (possible alignments)
 A (hypothesis alignments).




E.g.

Sure: 1-1 2-1 3-2 4-3 5-4 ...

Possible: 4-2 4-3 7-4 7-5

III. EXPERIMENTS



Corpus	Type	Name	No Sentences
Roman-English	Training	Naacl2003	48k
	Training	WMT2016 SETIMES	213k
	Testing	Naacl2003	248
English-Czech	Training	News commentary v.11	191432
	Testing	Marecek2008	2500
Dutch-English	Training	Europarl	2M
	Testing	Europarl	509
English-Italian	Training	Europarl	2M
	Testing	WAGS	6700

- ❑ Romanian-English testing set only includes sure alignments.
- ❑ English-Italian testing set includes only rare words.

Corpus	IBM2	IBM4	Best
Ro-En	30.7	30.4	IBM4
En-Cz	24.3	26.7	IBM2
Du-En	27.4	22.3	IBM4
En-It	68.6	80.6	IBM2

IBM2: Fast_align

IBM4: MGIZA++


III. EXPERIMENTS

Corpus	IBM2	IBM4	NWA-HMM
En-Cz	24.3	26.7	82.4

- 82.4 is still a very bad score.
- ✓ AER score has a decreasing tendency until 9th epoch

Other investigations:

- ✓ Maximum log-likelihood in Baum-Welch algorithm.
- ✓ The variation of non-negative transition elements $s(i-i')$ through epochs.




The AER scores for 10 first epochs
on English-Czech corpus

Maximum log-likelihood


BW Algorithm uses EM algorithm to find the maximum log-likelihood:

$$\theta^* = \arg \max_{\theta} P(f^J | \theta)$$




✓ Ascending trend through epochs

The variation of non-negative transition elements $s(i-i')$








✓ The sorter distance
the higher value

The mean of non-negative transition elements $s(i-i')$




- ✓ The sorter distance the higher value


The variation of non-negative transition elements $s(i-i')$



The variation of non-negative transition elements $s(i-i')$



Programing Tips and Tricks



A potential arithmetical issue during running BW-Alg
e.g. $x^{-200} \cdot x^{-200}$

1. Baum-Welch normalization
2. Log-space multiplication between two very small numbers dealing with normalization task

Programing Tips and Tricks

Baum-Welch normalization

Forward messages α and backward messages β
can get very small. e.g. x^{-200} .

Solution: Using the same normalization factor

$$Z(j) = \sum_i^I \alpha_i(j)$$

$$\hat{\alpha}_i(j) = \alpha_i(j)/Z(j)$$

$$\hat{\beta}_i(j) = \beta_i(j)/Z(j)$$

Programing Tips and Tricks

Log-space multiplication between two very small numbers dealing with normalization task

$$\log(\hat{x}_i) = \log(x_i) - \log(\sum_{j=0}^J (x_j)).$$

However $\sum_{j=0}^J (\log(x_j)) \neq \log(\sum_{j=0}^J (x_j)).$

$$\hat{x}_i = \frac{x_i}{\sum_{j=0}^J x_j}$$

$$\log(\sum_{j=0}^J (x_j)) = \log(x_0) + \log(1 + \sum_{j=1}^J \frac{x_j}{x_0})$$

$$= \log(x_0) + \log(1 + \sum_{j=1}^J (\exp^{\log(x_j) - \log(x_0)}))$$

$$= \log(a_0 \times b_0) + \log(1 + \sum_{j=1}^J (\exp^{\log(a_j \times b_j) - \log(a_0 \times b_0)}))$$

$$\log(x_i) = \log(a_i) + \log(b_i)$$

$$= \log(a_0) + \log(b_0) + \log(1 + \sum_{j=1}^J (\exp^{\log(a_j) + \log(b_j) - \log(a_0) - \log(b_0)}))$$

where $x_0 > x_1 > \dots > x_J$ are sorted in descending order.



IV. CONCLUSION

For enhancing translation performance:

- ✓ Believe that an improvement of this proposed method could be useful for further work in unsupervised neural alignment.
- ✓ Potentially to be integrated into well-know Attention Model.

Reasons for using neural alignment:

- ✓ Do not have much aligned corpus which requires expensively human resources.
- ✓ Have not yet found an unsupervised efficient automatic word alignment method which could obtain less than about 10% of error.
- ✓ The word alignment is still a promised model to enhance the translation achievements considerably.
- ✓ Neural network itself is really powerful for extracting special features on our data that plays a necessary role in each alignment



Thank you for your attention !